# HOWTO Write a UDenver Autopilot Plugin

The University of Denver autopilot is a highly extensible system, plugins can be written and built in to the system in short order meaning that you can add support for new devices or functionality quickly.

We'll create a plugin used for polling CPU load on the system and reporting it back to the ground station.

### Implementing the Driver

First, make a directory under src that will contain the files your plugin needs to operate, in this example we'll use linux and put in the files Linux.h and Linux.cc. Make sure you end your C++ files with the .cc extension!

This is our Linux.h file:

```
* Copyright 2014 Joseph Lewis <joseph@josephlewis.net>
 * This file is part of University of Denver Autopilot.
      UDenver Autopilot is free software: you can redistribute it
      and/or modify it under the terms of the GNU General Public
      License as published by the Free Software Foundation, either
      version 3 of the License, or (at your option) any later version.
      UDenver Autopilot is distributed in the hope that it will be useful,
      but WITHOUT ANY WARRANTY; without even the implied warranty of
      MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
      GNU General Public License for more details.
      You should have received a copy of the GNU General Public License
      along with UDenver Autopilot. If not, see <a href="http://www.gnu.org/licenses/">http://www.gnu.org/licenses/</a>>.
 *******************************
#ifndef LINUX_H
#define LINUX H
#include <atomic> // Used for atomic types
#include <mutex> // Used for singleton design.
#include "Driver.h" // All drivers implement this.
 * Provides an interface to the performance of Linux.
```

```
**/
   class Linux: public Driver {
   public:
       /**
       Returns the one allowed instance of this Driver
       static Linux* getInstance();
        * Returns the CPU utilization of the whole system.
       float getCpuUtilization() const { return cpu_utilization.load();}
       // We'll get to this in a bit
       virtual bool sendMavlinkMsg (mavlink_message_t* msg, int uasId, int sendRateHz, int
   private:
       static Linux* _instance;
       static std::mutex _instance_lock;
       Linux();
       virtual ~Linux();
       static void cpuInfo(Linux* instance); // this will be started in a new thread
       std::atomic<float> cpu_utilization; // store the value here
       bool isEnabled; // tells us if the plugin is enabled or not
   };
   #endif /* LINUX_H */
Section by section we have:
   * Copyright 2014 Joseph Lewis <joseph@josephlewis.net>
    * This file is part of University of Denver Autopilot.
          UDenver Autopilot is free software: you can redistribute it
          and/or modify it under the terms of the GNU General Public
          License as published by the Free Software Foundation, either
          version 3 of the License, or (at your option) any later version.
          UDenver Autopilot is distributed in the hope that it will be useful,
```

- \* but WITHOUT ANY WARRANTY; without even the implied warranty of
- \* MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
- \* GNU General Public License for more details.

\*

- \* You should have received a copy of the GNU General Public License
- \* along with UDenver Autopilot. If not, see <a href="http://www.gnu.org/licenses/">http://www.gnu.org/licenses/</a>.

This is the standard GNU GPL license, it ensures that all code you produce can be used with the UDenver Autopilot and alerts future developers to who owns the intellectual property.

Next up are the macros that ensure the file is only included once, be sure to name them appropriately for your file:

```
#ifndef LINUX_H
#define LINUX_H
```

Afterwards we have the includes, C++11 allows for built in atomic types and mutexes. We'll use atomics for anything that we want to store that is an int, float, int, bool, or char that needs to be modified in one thread and accessed in another. We'll use a mutex for locking the instance of this driver. Driver.h provides access to a collection of methods that provide logging mechanisms, configuration management, quick serial port configuration, and notification to shut down.

```
#include <atomic> // Used for atomic types
#include <mutex> // Used for singleton design.
#include "Driver.h" // All drivers implement this.
```

After we have our class definition:

```
class Linux: public Driver {
public:
```

We try to make all Drivers singetons, meaning only one can exist at a time in the application, and to get that one we just call getInstance() this makes sure you don't mistakenly create multiple instances and then get multiple messages back, and it also means you don't have to keep track of the object once the first getInstance() is called.

```
/**
Returns the one allowed instance of this Driver
**/
static Linux* getInstance();
```

Here we have a method that returns the CPU utilization (what we're measuring). Note that the method is const meaning it doesn't change any internal variables and it returns cpu\_utilization.load(), the .load() method of a std::atomic returns the value stored in that atomic. Usually it isn't needed, but trying to return a direct cpu\_utilization would be of type std::atomic<float> not float!

```
/**
 * Returns the CPU utilization of the whole system.
 **/
float getCpuUtilization() const { return cpu utilization.load();}
```

After that, we have a method that is overridden from Driver.h, by default you do not need to override this method if you are making a driver, just if you want the driver to report something back to QGroundControl.

This method is called once every time QGroundControl wants to send a message, it is passed in a blank mavlink\_message\_t, the Id of the helicopter, the number of packets sent to QGroundControl per second and the total number of messages sent so far. These last two can be used to limit the number of messages you send to a certain rate.

Our private variables are fairly straightforward:

- \_instance holds the created instance of this object
- $\bullet$  \_instance\_lock locks the instance so we don't inadvertantly create multiple
- cpu utilization stores the current utilization
- isEnabled stores the configuration value of whether or not this driver is enabled. We want to store it because looking it up every time is expensive.

```
private:
    static Linux* _instance;
    static std::mutex _instance_lock;
    std::atomic<float> cpu_utilization; // store the value here
    bool isEnabled; // tells us if the plugin is enabled or not
```

After those comes the standard constructor and destructor:

```
Linux();
virtual ~Linux();
```

Finally we have the method that does the fun stuff. It is what is started in a new thread to do the reading. We pass a pointer to instance rather than having it use getInstance() inside the cpuInfo method because there is no gaurantee that getInstance() will return the actual instance immediately when this thread begins, it could instead hang the program!

static void cpuInfo(Linux\* instance); // this will be started in a new thread

Now let's look at our C++ definition file:

```
* Copyright 2014 Joseph Lewis <joseph@josephlewis.net>
 * This file is part of University of Denver Autopilot.
      UDenver Autopilot is free software: you can redistribute it
      and/or modify it under the terms of the GNU General Public
      License as published by the Free Software Foundation, either
      version 3 of the License, or (at your option) any later version.
      UDenver Autopilot is distributed in the hope that it will be useful,
      but WITHOUT ANY WARRANTY; without even the implied warranty of
      MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
      GNU General Public License for more details.
      You should have received a copy of the GNU General Public License
      along with UDenver Autopilot. If not, see <a href="http://www.gnu.org/licenses/">http://www.gnu.org/licenses/</a>.
 *******************************
#include "Linux.h"
#include <boost/thread.hpp>
#include <iostream>
#include <fstream>
#include "RateLimiter.h"
Linux* Linux::_instance = NULL;
std::mutex Linux:: instance lock;
Linux* Linux::getInstance()
```

```
std::lock_guard<std::mutex> lock(_instance_lock);
    if (_instance == NULL)
    {
        _instance = new Linux;
    return _instance;
}
Linux::Linux()
:Driver("Linux CPU Info", "linux_cpu_info")
{
    isEnabled = configGetb("enabled", true);
    // If the system wants to halt, don't start running.
    if(terminateRequested())
        return;
    }
    // If the user has disabled this component, don't start running
    if(! isEnabled)
        return;
    // Tell the user we made it up
    warning() << "starting CPU Information System";</pre>
    // Start our processing thread.
    boost::thread(std::bind(Linux::cpuInfo, this));
}
Linux::~Linux() {
}
void Linux::cpuInfo(Linux* instance)
    RateLimiter rl(1);
    std::ifstream myfile;
    myfile.open ("/proc/loadavg");
    float util = 0;
    while(!instance->terminateRequested())
```

```
rl.wait(); // suspends until we're ready to go
        myfile >> util;
        myfile.seekg(0); // go to the beginning of the file
        instance->cpu_utilization = util; // set our utilization
        instance->debug() << "Got Load of: " << util;</pre>
        rl.finishedCriticalSection(); // call to yield
    }
    myfile.close();
}
bool Linux::sendMavlinkMsg(mavlink message t* msg, int uasId, int sendRateHz, int msgNu
{
    if(! isEnabled) return false;
    if(msgNumber % sendRateHz == 0) // do this once a second.
    {
        debug() << "Sending CPU Utilization";</pre>
        mavlink_msg_udenver_cpu_usage_pack(uasId, 40, msg, getCpuUtilization());
        return true;
    return false;
};
```

Again, let's go section by section ignoring the includes and copyright this time. First, we must set defaults for the mutex and instance, this ensures that getInstance() will work as anticipated:

```
Linux* Linux::_instance = NULL;
std::mutex Linux::_instance_lock;
```

Next, we write the getInstance() method, the first time it is called the lock gets locked, \_instance is NULL and we create a new object then return it. Each time after we see \_instance is no longer NULL and we return the instance.

```
Linux* Linux::getInstance()
{
    std::lock_guard<std::mutex> lock(_instance_lock);
    if (_instance == NULL)
    {
        _instance = new Linux;
    }
}
```

```
return _instance;
}
```

The constructor comes after, we must instantiate Driver with two things:

- 1. A human readable name for this Driver used in logging
- 2. A valid XML tag name used for storing this Driver's configuration.

Immediately after staring we set our variables, and check if terminateRequested is true, this is a method provided by Driver that alerts all components of the system that the Helicopter wants to quit. Once they check it they should prepare themselves for termination by saving all needed data and then exiting.

After that we quit if the user had explicitly disabled this component in config.xml.

Finally we emit a message telling the user that we are going to start up, and call the cpuInfo method in a new thread passing it this as the instance param.

```
Linux::Linux()
:Driver("Linux CPU Info", "linux_cpu_info")
{
    isEnabled = configGetb("enabled", true);
    // If the system wants to halt, don't start running.
    if(terminateRequested())
    {
        return;
    }
    // If the user has disabled this component, don't start running
    if(! isEnabled)
        return;
    }
    // Tell the user we made it up
    warning() << "starting CPU Information System";</pre>
    // Start our processing thread.
    boost::thread(std::bind(Linux::cpuInfo, this));
}
```

Our destructor is empty for now because the object is never destroyed.

```
Linux::~Linux() {
}
```

The cpuInfo method is fairly simple. Most of the time Drivers are interacting with a serial port and wait when the read and write to that, however Linux reads from the local machine, therefore we must rate limit it to a reasonable speed, in this case 1Hz. This limiting keeps it from doing this thousands of times a second and maxing out the CPU.

After we create our RateLimiter we open the file, specified by the proc man page this file is a special one provided by the Linux kernel, it is updated every five seconds with information about the CPU usage of the machine.

Once the file is open, we loop until the autopilot wants to shut down updating as we go along, and finally close the file.

```
void Linux::cpuInfo(Linux* instance)
{
    RateLimiter rl(1);
    std::ifstream myfile;
    myfile.open ("/proc/loadavg");
    float util = 0;
    while(!instance->terminateRequested())
        rl.wait(); // suspends until we're ready to go
        myfile >> util;
        myfile.seekg(0); // go to the beginning of the file
        instance->cpu_utilization = util; // set our utilization
        instance->debug() << "Got Load of: " << util;</pre>
        rl.finishedCriticalSection(); // call to yield
    }
    myfile.close();
}
```

The last method is fairly straightfoward. When QGroundControl wants to emit more messages, it calls this on every Driver. If Linux is not enabled, then we return false meaning we don't have a message to send, otherwise we'll try to send once a second (the message number modulo the send rate).

At that time, we'll send a debug message, pack the message using a method automatically provided by Mavlink once we tell it about our message, and return

true. QGCLink will handle the rest and your message will begin to be parsed by QGroundControl on the other side.

```
bool Linux::sendMavlinkMsg(mavlink_message_t* msg, int uasId, int sendRateHz, int msgNum
{
    if(! isEnabled) return false;

    if(msgNumber % sendRateHz == 0) // do this once a second.
    {
        debug() << "Sending CPU Utilization";
        mavlink_msg_udenver_cpu_usage_pack(uasId, 40, msg, getCpuUtilization());
        return true;
    }
    return false;
};</pre>
```

#### Integrating With the Autopilot

Once we have this simple skeleton of a plugin, we can add initializing it to the MainApp class which is responsible for starting all extensions. At the top, include Linux.h and in the MainApp::run() method, call getInstance() after all the other plugins have done the same:

```
...
message() << "Setting up altimiter";
MdlAltimiter::getInstance();
message() << "Setting up Linux CPU Reader";
Linux::getInstance();
...</pre>
```

It is polite to leave a message before your getInstance() in the event your plugin hangs the system the user will be able to tell where it happened.

The RateLimiter class suspends a thread until it is time to wake up and do work, calling the finishedCriticalSection() allows the thread to give up control of its remaining time. This ensures the process runs smoothly and is kind to all other threads on the system.

# Sending Data to QGroundControl

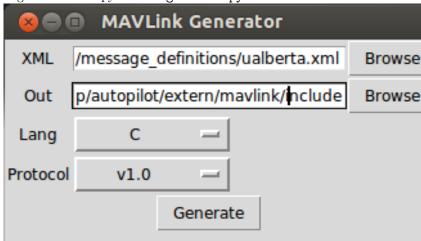
Data is ferried between QGroundControl and the autopilot using MavLink packets.

MavLink auto-generates itself from the ualberta.xml file found in extern/mavlink/message\_definitions:

Add the following under the messages section, ensuring the ID used is unique and less than 255:

```
<message id="230" name="UDENVER_CPU_USAGE">
     <description>Autopilot System CPU Usage</description>
     <field type="float" name="cpu_usage"/>
</message>
```

Then generate the new ualberta xml using the included python mavgenerate.py



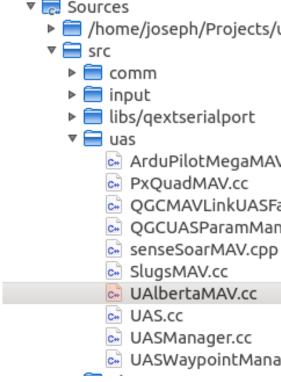
script in to extern/mavlink/include

This generates new headers in the extern/mavlink/include directory. Also copy extern/mavlink to the same directory as QGroundControl so it gets the updated messages as well.

#### Integrating with QGroundControl

Be sure to copy the mavlink directory to the same parent folder as QGround-Control first.

- 1. Open QTCreator
- 2. Open the QGroundControl project



- 3. Open the UAlbertaMav.cc file under Sources > Src > uas
- 4. Scroll to the end of the large switch block that begins on line 40.
- 5. Before the default add in the switch for your new message:

```
case MAVLINK_MSG_ID_UDENVER_CPU_USAGE:
{
    // saves the CPU usage from mavlink
    mavlink_udenver_cpu_usage_t cpu;

    // converts the generic message we got to the CPU one we understand
    mavlink_msg_udenver_cpu_usage_decode(&message, &cpu);

    // emit a "CPU usage" event, which allows us to graph the usage
    // we use the current UNIX time as the timestamp, although you may
    // want to get this from the system instead, left as an exercise
    // for the reader.
    emit valueChanged(uasId, "CPU usage", "amt", cpu.cpu_usage, getUnixTime());
    break;
}
```

6. Save, compile, and run QGroundControl

# Running

Now we can build the autopilot and your plugin will be added in! If it is not, make sure your C++ file has the .cc extension.

Now run the autopilot, if you don't see your extension starting, check to make sure that your <enabled> tag in the config.xml file is true. You should see a message like this when your extension begins:

```
Debug: Linux CPU Info: starting CPU Information System
```

In config.xml you should see a definiton like this:

Let's set debug to true and enabled to true as well, ignoring the read\_style for now. When debug is true, it allows the printing of debugging messages from the module and also turns all trace messages into logged messages.

